


Adversarial Search I

The Core Idea

- Let's say you are playing a game
 - Used to be playing a 1 player game
 - (Even if >1 player, agent's "search" didn't consider other player(s) actions)
- What if we know the other player(s) move(s)?
 - Can we model their goals?
 - Can we predict what actions they will take?
 - If so, can we make better choices?

Updating Some Terms


- Utility function $u(s) \rightarrow u(s, p)$  Only works on terminal states
- Transition function $t(s, a) \rightarrow t(s, p, a)$
- Previously talked about goal function $g(s)$
 - More general term: terminal test function *terminal* – $test(s)$
- Game is a 6 tuple now:
 - Initial state s_0
 - Set of players
 - Actions available to each player in each state (assume same actions for now)
 - *terminal* – $test(s)$
 - $u(s, p)$
 - $t(s, p, a)$

How can We Make Better Choices?

Turn 1
 $p_1 \rightarrow p_2 \rightarrow p_3 \rightarrow \dots \rightarrow p_n \rightarrow p_1 \rightarrow p_2 \rightarrow \dots \rightarrow p_n \rightarrow \dots$
Turn 2
...


- Lets assume game has **known order of turns**
 - Each player goes **one at a time**
 - **Static world** (world pauses while each agent thinks)

p_1 's perspective



What to Do?

- Can expand this tree all the way down to **terminal states**
 - Pick the path that ends with a terminal state that's good for us?
- Why doesn't this work?
 - Practical problems:
 - Tree is **massive**!
 - Takes **too long to expand** the whole thing
 - Theoretical problems:
 - We only know utility value of terminal states:
 - How do we determine utility values of non-terminal states?
 - Combinatorial ways of combinations (b/c of lots of players)
- Need:
 - **Faster** tree algorithms
 - Polynomial time ways of combining other player's move(s)



Simplify (for now)

- Consider a **2-player game**
 - agent1 vs agent 2 (competitive/adversarial game)
 - Zero sum game** (a major simplifying assumption)
 - My loss is your gain (and vice versa)
 - "true utilities" always sum to same constant (may not be zero!)
- Adversary wants to force me to go to bad states (**for me**)
 I want to force adversary in bad states (**for them**)
- With these two assumptions:
 - Tree is much smaller
 - Don't need to worry about combining multiple adversaries together
 - What's bad for me is good for you
 - What's good for me is bad for you
 - Nonterminal states utility value = function of children utility values

$$u(s, p) = \begin{cases} \text{terminal} - \text{test}(s) \cap p == \text{me} & \text{How good my (nonterminal) state is} \\ \text{terminal} - \text{test}(s) \cap p! = \text{me} & \text{How good adversary's (nonterminal) state is} \\ !\text{terminal} - \text{test}(s) \cap p == \text{me} & \max_a u(t(s, p, a), \text{me}) \\ !\text{terminal} - \text{test}(s) \cap p! = \text{me} & \min_a u(t(s, p, a), \text{me}) \end{cases}$$

The best state for me in the future!

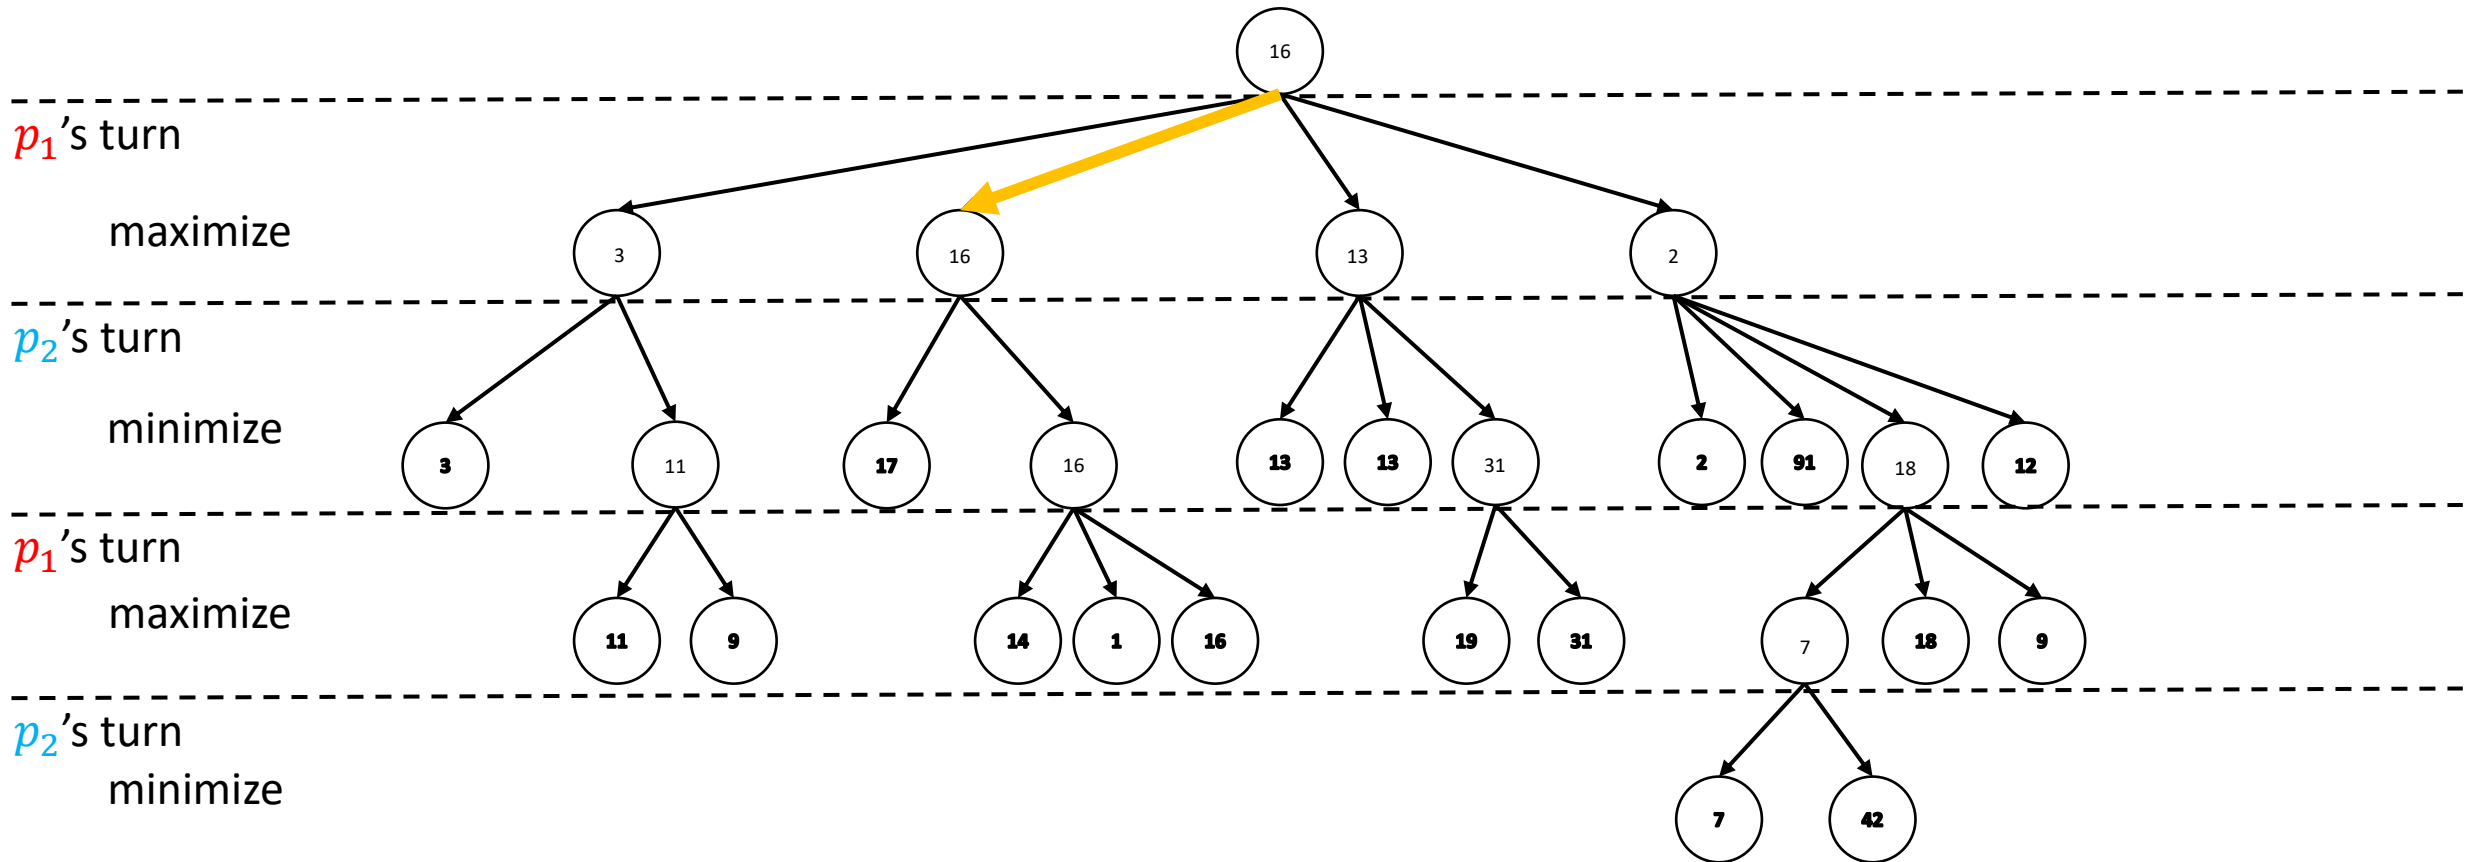
$u(s)$
 $u(s)$

The worst state for me in the future

How good adversary's (nonterminal) state is

The Minimax Algorithm

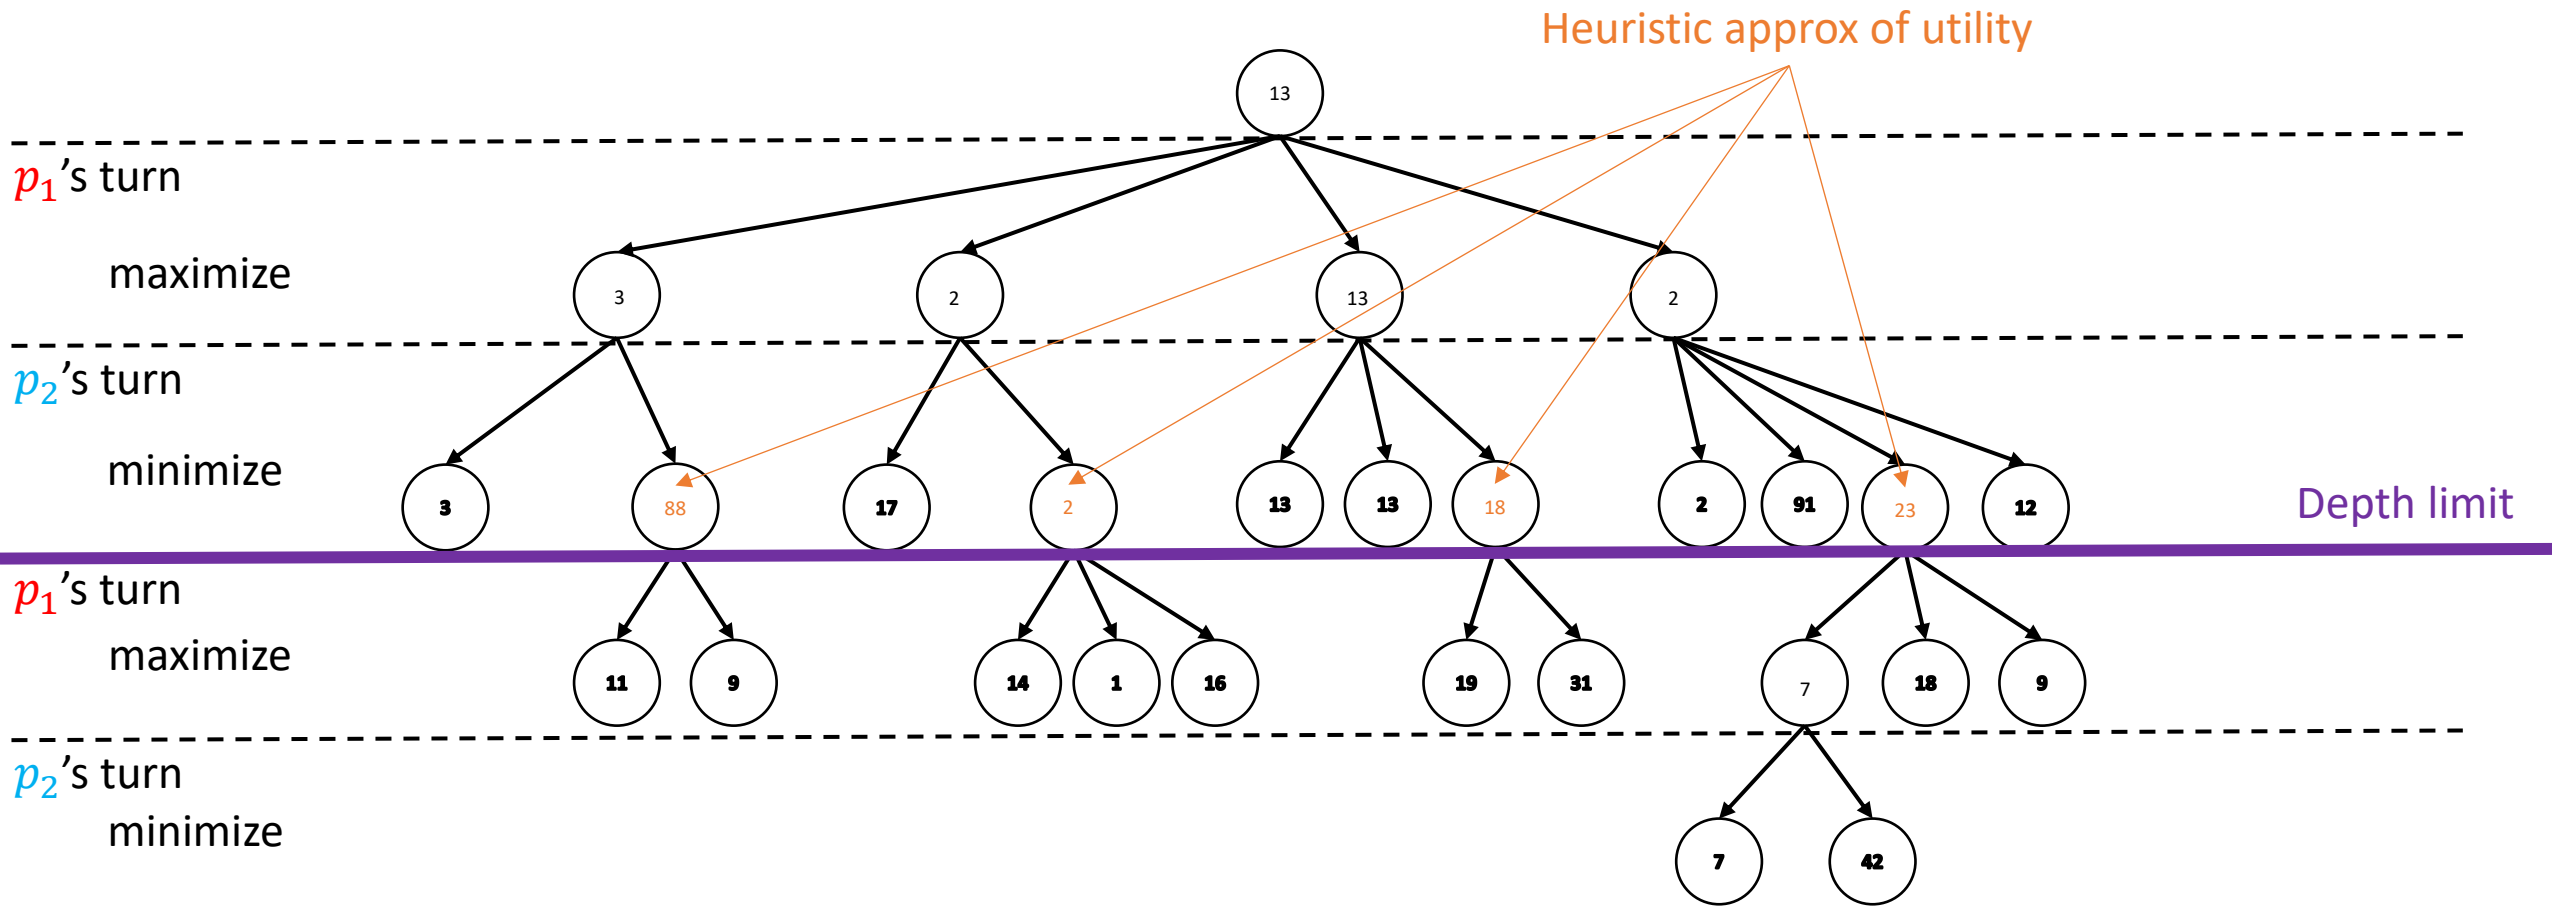
- Expand the tree (dfs is common)
- Once you get to terminal states (leaf nodes):
 - Go back to parent nodes and assign utilities
- At root:
 - Pick action which leads to the best child state



Is this a Good Model?

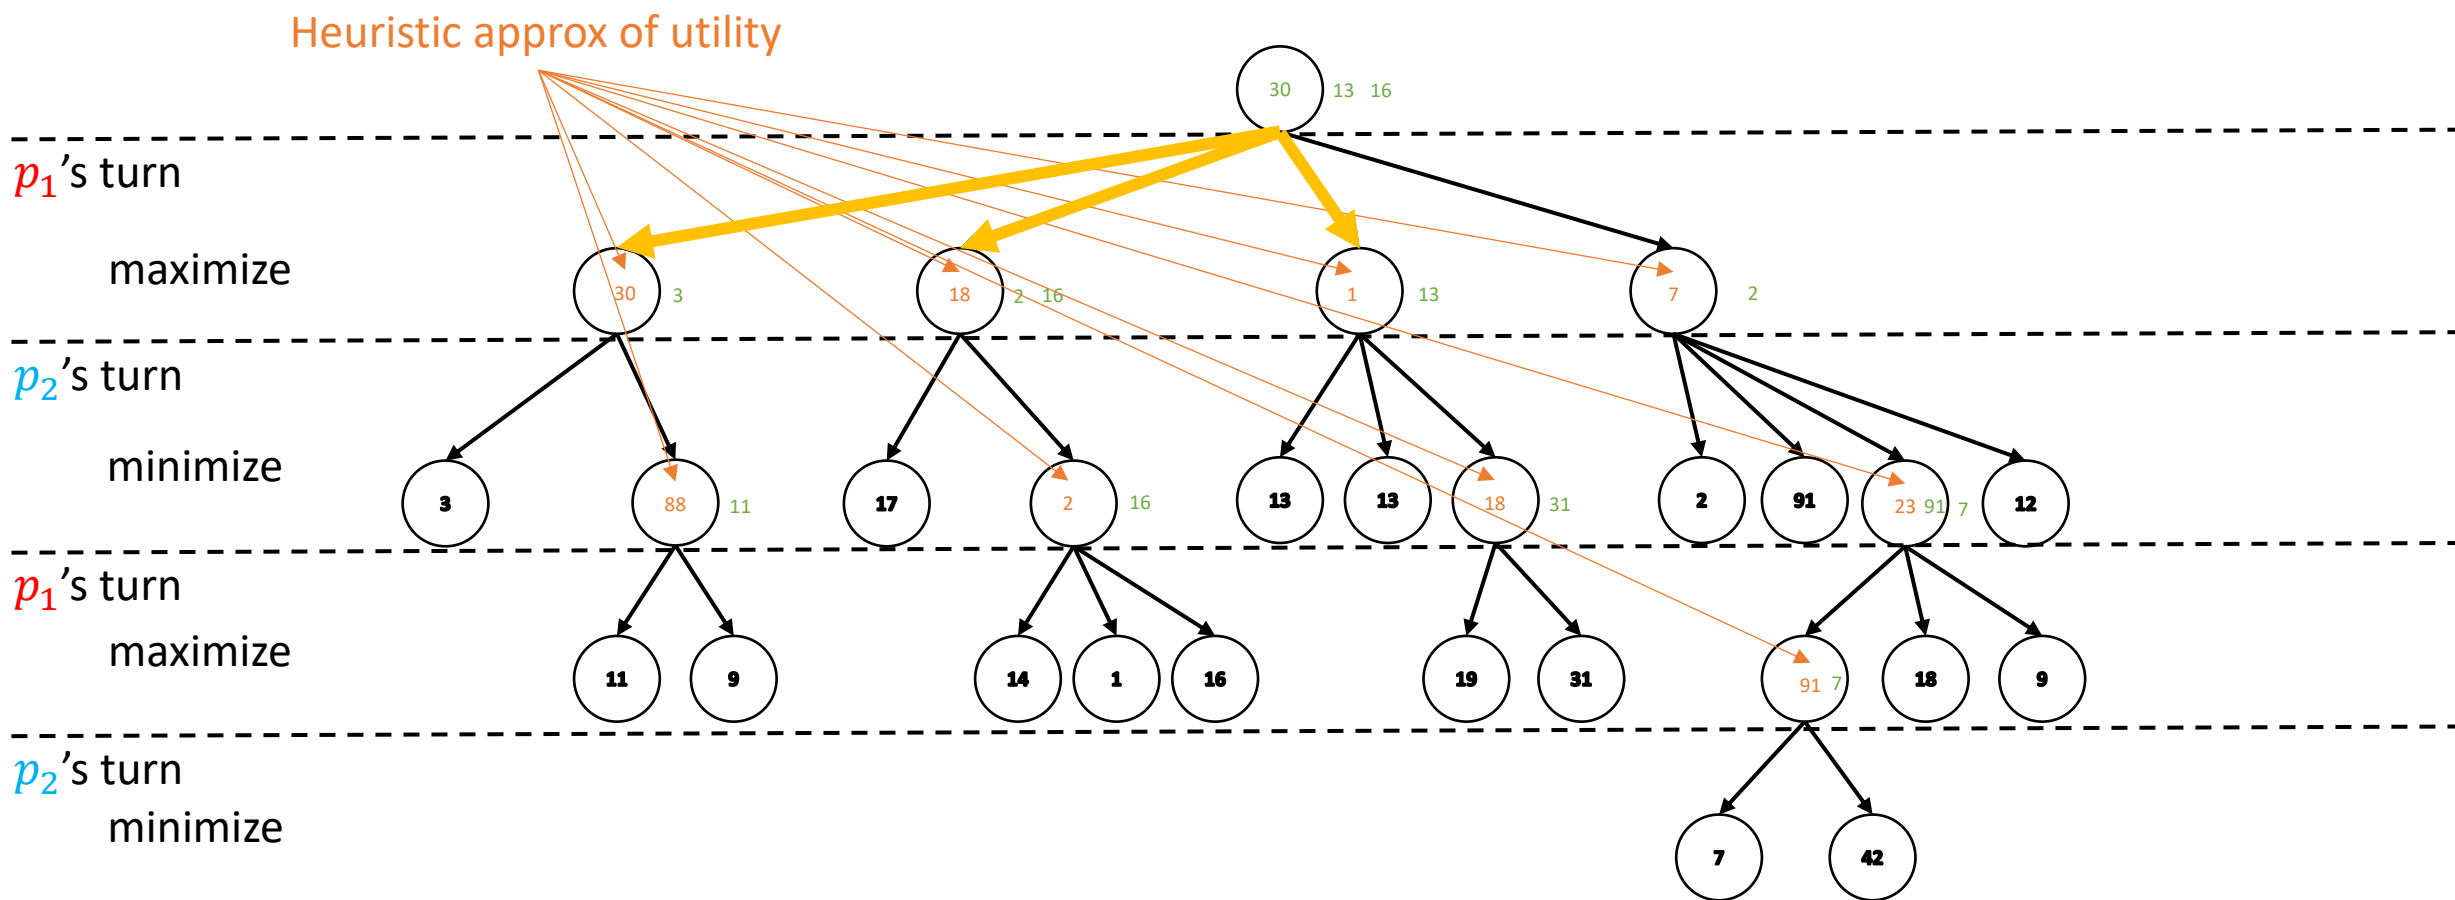
- Assumes adversary always makes the “best” choice
 - Humans are notorious for not doing this
 - How you beat chess bots!
 - Hint: don’t play chess with them (you will lose)
 - Instead:
 - Take advantage of time limits: make them think
 - If they run out of time they might play sub-optimally
- Major problem:
 - Even these trees can be too big
 - Idea: don’t go all the way to terminal states!
 - Ex. only plan “3 moves ahead” or something
 - Need a stand in for utility values of leaf nodes in the tree
 - heuristics

Depth-Thresholded Minimax



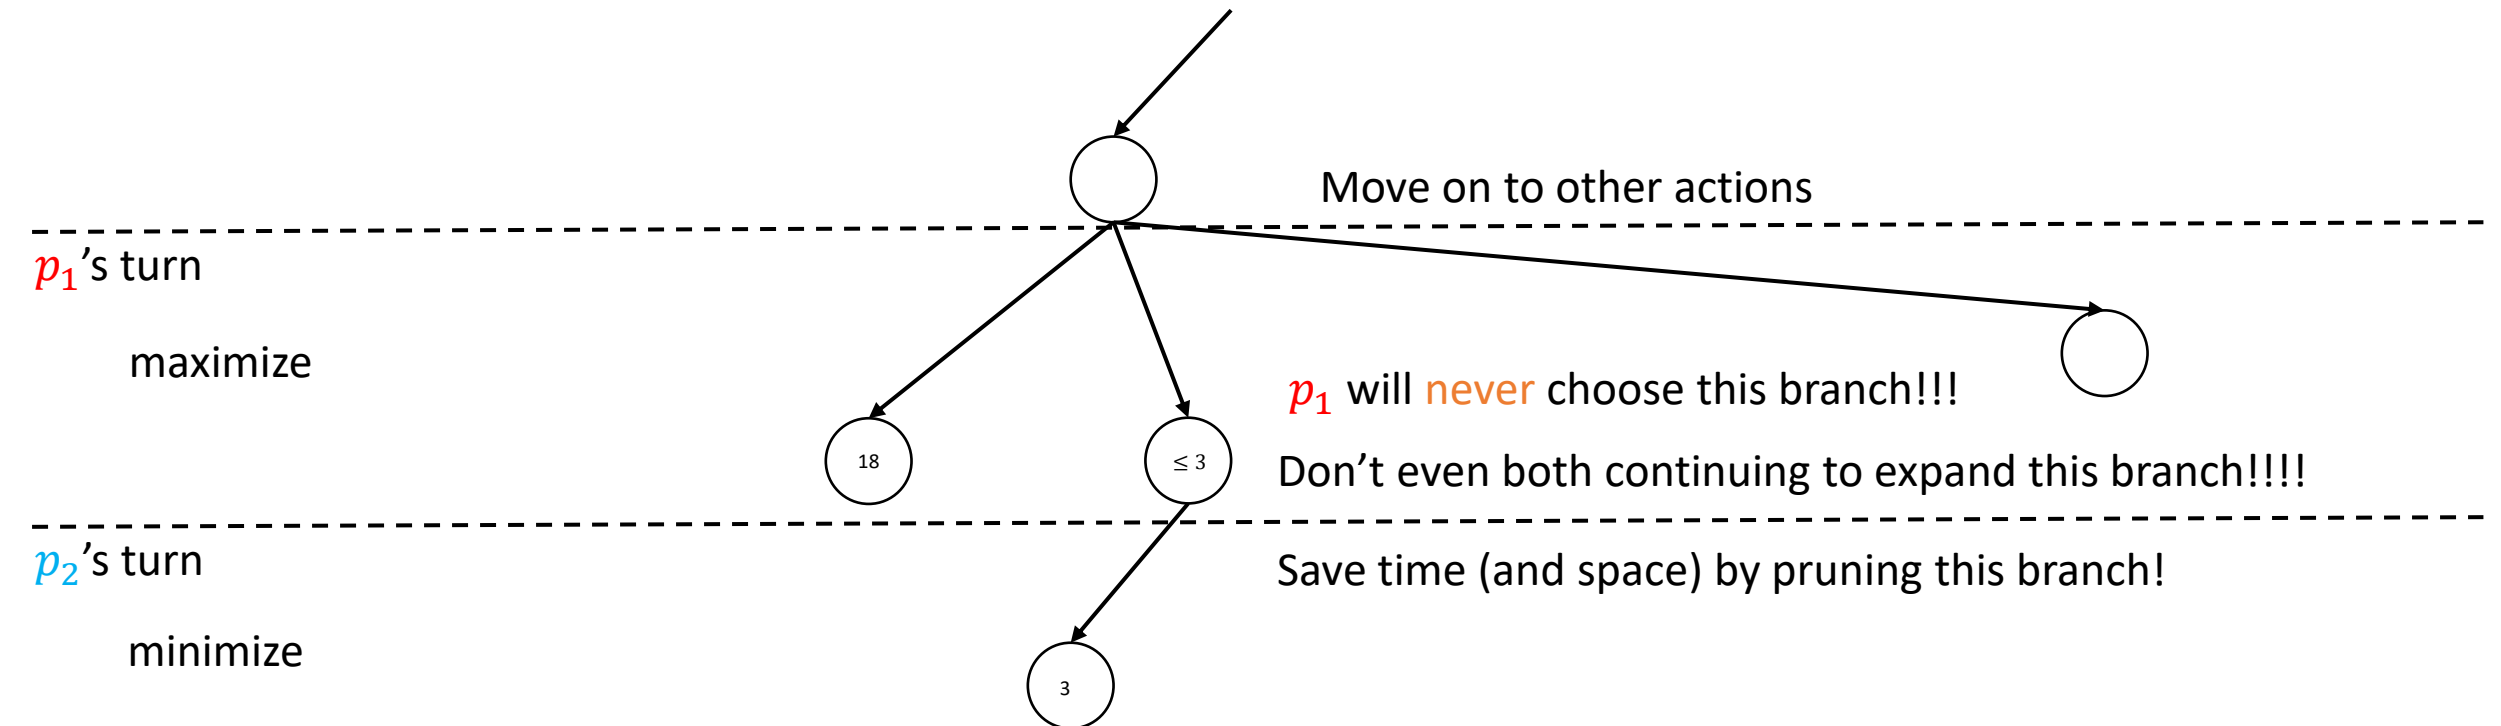
Iterative Deepening

- Expand the tree using **BFS**
 - Use heuristics to determine best move at a level
 - If you have more time/resources: do next level
- Whenever resources run out: return best choice so far



Problems

- Minimax is inefficient
 - Inefficiency is subtle
 - Iterative deepening / depth-thresholding don't solve it



Alpha-Beta Pruning

- How do we detect scenarios like this?

- Need to keep track of “best choice” for each player

Max player is guaranteed at least this score

- α = min score of the maximizing player (me) so far

Min player is guaranteed at most this score


- β = max score of the minimizing player (adversary) so far

- Every node gets its own pair of values (α, β)

- Child nodes inherit values of parents

- Child can change it's own values of (α, β) values

- Alpha-Beta pruning is Minimax w/ “early stopping”




When min player:

If I encounter child state $< \alpha$:
max player will never get here
stop expanding this branch!

When max player:

If I encounter child state $> \beta$:
min player will never get here
stop expanding this branch!



Alpha-Beta Pruning

```
function alphabeta( $s, \alpha, \beta, p$ ):
  if terminal – test( $s$ ) then:
    return  $u(s)$ 
```

```
  if is – maximizing – player( $p$ ):
```

```
     $v = -\infty$ 
```

```
    for each child state  $s'$  of  $s$  do:
```

```
       $v = \max(v, \text{alphabeta}(s', \alpha, \beta, \text{other} - \text{player}(p)))$ 
```

```
      if  $v > \beta$  then:
```

```
        return  $v$ 
```

```
       $\alpha = \max(\alpha, v)$ 
```

```
    return  $v$ 
```

```
  else:
```

// minimizing player

```
     $v = +\infty$ 
```

```
    for each child state  $s'$  of  $s$  do:
```



```
       $v = \min(v, \text{alphabeta}(s', \alpha, \beta, \text{other} - \text{player}(p)))$ 
```

```
      if  $v < \alpha$  then:
```

```
        return  $v$ 
```

```
       $\beta = \min(\beta, v)$ 
```

```
    return  $v$ 
```



Alpha-Beta Pruning

MAX player (you)

$v = -\infty$

for each child state s' of s do:

$v = \max(v, \text{alphabeta}(s', \alpha, \beta, \text{other} - \text{player}(p)))$

if $v > \beta$ then:

return v

$\alpha = \max(\alpha, v)$

return v

MIN player (adversary)

$v = +\infty$

for each child state s' of s do:

$v = \min(v, \text{alphabeta}(s', \alpha, \beta, \text{other} - \text{player}(p)))$

if $v < \alpha$ then:

return v

$\beta = \min(\beta, v)$

return v

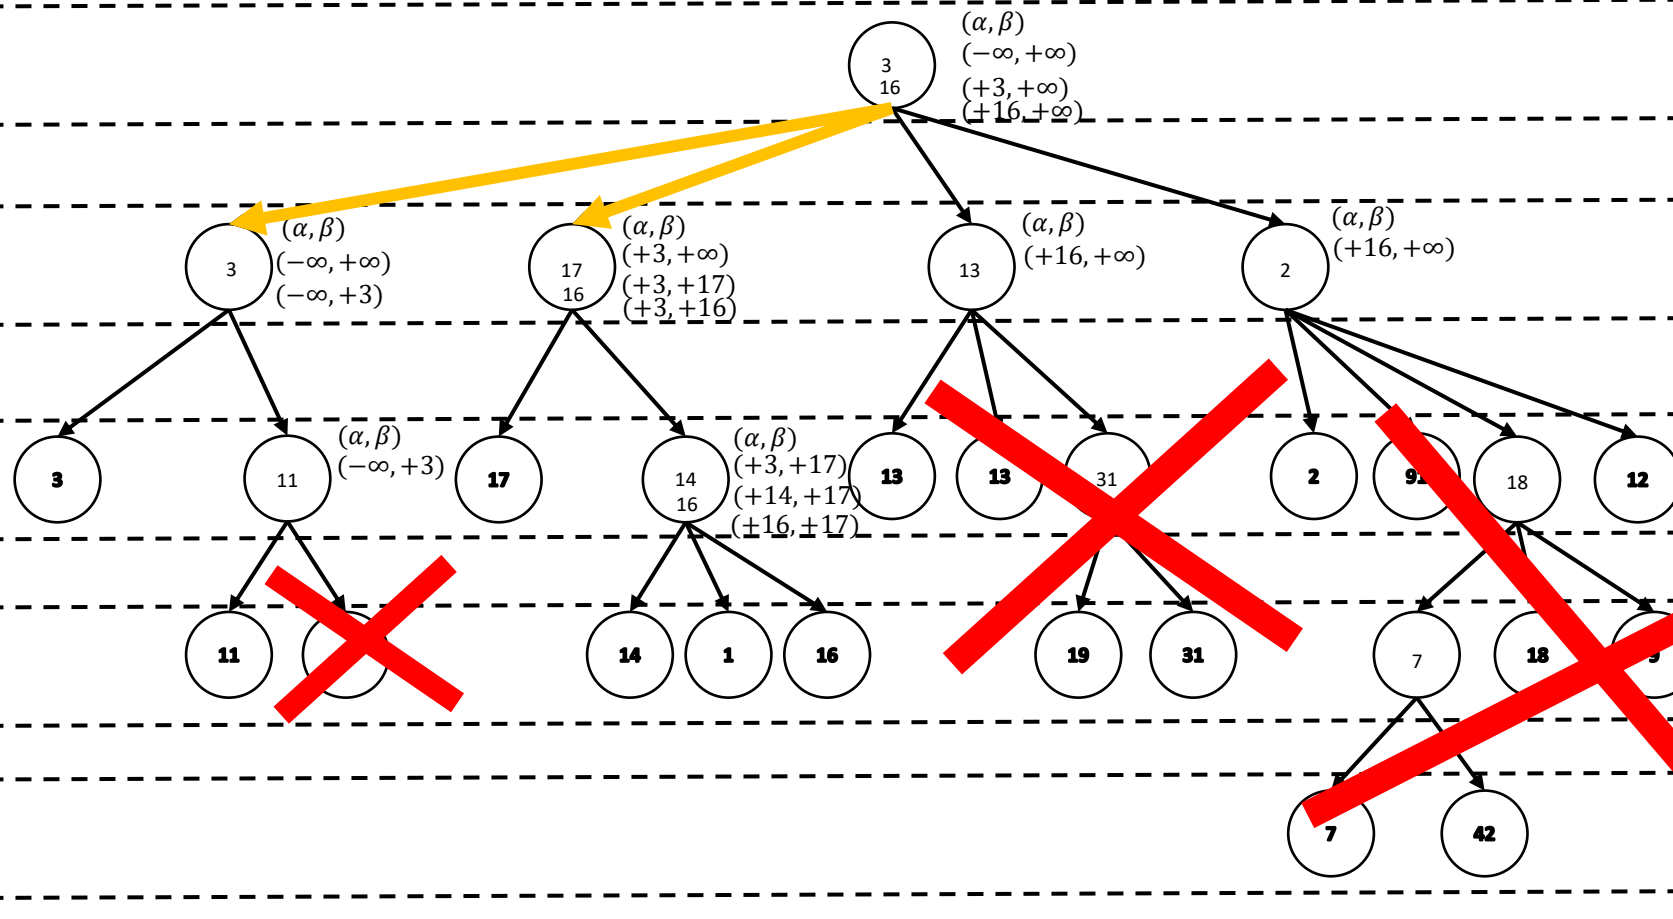
maximize

minimize

maximize

minimize

maximize



Variants of Alpha-Beta Pruning

- Even though we prune: Alpha-Beta still examines all the way to terminal states
- Solved this with Minimax:
 - Depth threshold ← Most common!
 - Iterative deepening
- These solutions work with Alpha-Beta too!

Problems with Alpha-Beta Pruning

- Even with Depth-Thresholding/Iterative Deepening:
 - Order in which child states are enumerated matters!
 - Pruning only occurs when we know a better option **already exists!**
 - What if we see better options last?
 - **Never prune!**
 - Without pruning, Alpha-Beta is just **Minimax!**
- How can we speed this up?
 - Transposition table:
 - Cache utility value for states we've seen
 - When we encounter the same state in the future, no need to expand!
 - How many states should we cache?
 - How can we encourage Alpha-Beta to prune?
 - More heuristics!
 - Impose an order on child enumeration
 - Children we think are better choices should come first in the order!

There are other ideas
Can implement multiple!

← **Super popular**